

# CallSphere Assignment: Multi-Agent Logistics & Delivery Platform (OpenAI-Only, Chat + Voice)

## 1) Objective

Design and implement an end-to-end **logistics & last-mile delivery** platform with:

- A **Customer dashboard** that supports **both** a **chat agent** and a **browser voice agent**.
- An **Admin/Ops dashboard** for real-time monitoring, issue handling, escalation workflows, and **editable metric configuration**.

The system must be production-minded (clean architecture, RBAC, auditability) while remaining feasible as an assignment prototype.

---

## 2) Required Tech Stack

- **Backend:** NestJS + Prisma + PostgreSQL
- **Frontend:** React + TypeScript + TailwindCSS
- **Realtime:** WebSockets (NestJS Gateway + frontend WS client)
- **AI (OpenAI only):**
- **OpenAI Agents** for multi-agent orchestration (router + specialist agents)
- **OpenAI Realtime API** for browser **voice** (STT/TTS) and streaming responses

**Constraints:** - Use **OpenAI infrastructure only** for AI/voice (no other LLM providers, no self-hosted models). - The product must support **both** a chat agent **and** a voice agent (not one or the other).

---

## 3) Users, Roles, and RBAC

Implement role-based access control (RBAC) across REST and WebSockets.

### Roles

- **customer**
  - Track own shipments (or public tracking by trackingNumber if you choose).
  - Create issues and request delivery changes for own shipments.
- **driver**
  - View assigned route(s) and stops.
  - Update stop status (completed/failed) for assigned shipments.
- **dispatcher**
  - View/update shipments and issues within assigned region.
  - Manage route adjustments (within policy).
  - Read metrics.
- **manager**
  - All dispatcher capabilities plus escalation acknowledgments.
  - Read metrics; optionally edit certain thresholds if allowed.
- **admin**
  - Full access.

- Can edit **metric definitions, thresholds, dashboard layouts, and escalation ladders**.

## RBAC Requirements

- Enforce RBAC in **NestJS Guards** on all REST endpoints.
  - Enforce RBAC for **WebSocket channels** (only authorized users can subscribe to route/issue/escalation feeds).
  - Frontend must hide/disable UI controls based on role.
- 

## 4) Product Scope

### 4.1 Customer Capabilities (Chat + Voice)

1. **Track Shipment**
2. "Where is my package ABC123?"
3. Returns status, last scan, and ETA.
4. **Report Delivery Issue**
5. Damaged, missing, wrong address, missed delivery, delay, other.
6. Creates a DeliveryIssue and notifies Ops in real time.
7. **Request Delivery Change**
8. Reschedule delivery window.
9. Update delivery instructions.
10. Change address (only when allowed by policy).

### 4.2 Ops/Admin Capabilities

1. **Routes & Stops**
  2. View route list and route detail (stops, statuses).
  3. Driver-level view (what a driver sees).
  4. **Issues & Escalations**
  5. Issue queue (filter by severity/status/type/region).
  6. Escalation ladder execution and acknowledgment.
  7. **Metrics Dashboard + Editing**
  8. Overview KPIs (on-time rate, first-attempt success, open issues, SLA-risk count).
  9. Admin can **edit metric definitions, targets, and thresholds**.
  10. Store metric snapshots for performance.
- 

## 5) Architecture

### 5.1 Backend (NestJS)

**Modules (suggested):** - auth, users - shipments, shipment-scans - routes, route-stops - delivery-issues - escalations, acknowledgments - metrics (compute + snapshots) - dashboard-config - agent-sessions (chat/voice sessions) - settings (policy controls) - ws (WebSocket gateway)

**Backend responsibilities:** - Owns the database and domain logic. - Exposes REST APIs (frontend + OpenAI tool calls). - Emits WebSocket events for live updates. - Performs background jobs (SLA risk scanning, metric snapshots).

## 5.2 AI Layer (OpenAI Agents + Realtime)

All agent logic runs through OpenAI: - Define a **router agent** and multiple **specialist agents**. - Provide tools so agents can safely call backend APIs.

**Chat path:** - User message → OpenAI Agent session → tool calls → response text.

**Voice path:** - Browser audio → OpenAI Realtime (STT) → agent routing/tools → response text + OpenAI TTS playback.

## 5.3 Realtime (WebSockets)

Backend publishes real-time events: - Shipment scan updates - Issue created/updated - Escalation triggered/advanced/acknowledged - Optional: metric snapshot updates

**Suggested channels:** - shipments:<trackingNumber> - routes:<routeCode> - issues - escalations - metrics:overview

---

# 6) Core Data Model (Domain)

Implement at least these tables (Prisma models) with appropriate indexes.

## 6.1 User

- id, name, email, role, isActive, createdAt, updatedAt

## 6.2 Shipment

- id
- trackingNumber (unique, indexed)
- orderId (optional)
- customerId (FK → User)
- fromAddress, toAddress
- currentStatus (enum)
- serviceLevel (enum)
- promisedDeliveryDate
- lastScanAt, lastScanLocation
- isVip (bool)
- slaRiskScore (float 0-1)
- createdAt, updatedAt

## 6.3 ShipmentScan

- id, shipmentId (indexed), scanType, location, timestamp (indexed), notes

#### **6.4 Vehicle**

- `id`, `vehicleCode` (unique), `capacityVolume`, `capacityWeight`, `homeBase`

#### **6.5 Driver**

- `id`, `driverCode` (unique), `userId`, `assignedVehicleId`, `homeBase`

#### **6.6 Route**

- `id`, `routeCode` (unique), `date` (indexed), `driverId`, `vehicleId`, `region`

#### **6.7 RouteStop**

- `id`, `routeId` (indexed), `shipmentId` (indexed), `sequenceNumber`, `plannedEta`, `actualArrival`, `status`

#### **6.8 DeliveryIssue**

- `id`, `shipmentId` (indexed), `reportedById`
- `issueType`, `description`
- `aiSeverityScore` (float 0-1)
- `status`, `resolutionNotes`
- `createdAt`, `updatedAt`

#### **6.9 EscalationContact**

- `id`, `userId`, `position`, `contactType`, `timeoutSeconds`, `isActive`

#### **6.10 EscalationLog**

- `id`, `shipmentId`, `deliveryIssueId` (optional)
- `contactId`, `attemptNumber`
- `eventType`, `payload` (JSON)
- `ackReceived`, `ackMethod`, `acknowledgedAt`, `createdAt`

#### **6.11 Acknowledgment**

- `id`, `shipmentId`, `deliveryIssueId` (optional), `userId`, `method`, `notes`, `createdAt`

---

## **7) AI Agents (OpenAI)**

Define the following agents (names are suggestions):

### **1. LogisticsRouterAgent**

2. Classifies intent and routes to a specialist agent.

### **3. ShipmentTrackingAgent**

4. Resolves tracking number, fetches shipment + latest scan, returns status & ETA.

#### 5. **DeliveryIssueAgent**

6. Collects issue details, classifies issueType, assigns aiSeverityScore, creates issue.

#### 7. **DeliveryChangeAgent**

8. Validates and applies delivery changes based on shipment status/policies.

#### 9. **LogisticsEscalationAgent**

10. Triggers escalation for high severity, VIP, or SLA-risk shipments.

11. Advances ladder until ACK is recorded.

#### 12. **LogisticsAnalyticsAgent**

13. Answers ops questions using metrics endpoints/snapshots.

**Tooling requirement:** - Agents must use backend tools to read/write data (no hallucinated updates).

---

## 8) Dashboards (Both Modalities Required)

### 8.1 Customer Dashboard (Chat + Voice)

Build a customer page that includes: - Tracking input + shipment list. - Shipment detail panel (timeline of scans, current status, ETA, open issues). - **Chat + Voice Agent widget:** - Text input + send. - Mic start/stop. - Live transcription for voice. - Unified conversation history across chat and voice.

Must support: - Track shipment - Create issue - Request change

### 8.2 Admin/Ops Dashboard (Metrics + Editing)

Build an internal dashboard with: - **Overview KPIs** (cards + trends) - **Shipments** (filters, detail) - **Routes** (list + route detail) - **Issues** (queue, severity, SLA risk) - **Escalations** (history + acknowledgment actions) - **Metrics Admin** (admin-only): - Edit metric definitions, targets, thresholds - Toggle visibility on dashboards - Manage dashboard layouts (role defaults)

---

## 9) Database Schema Additions (Single Place)

All dashboard + agent-specific schema extensions must live only here.

### 9.1 **AgentSession**

- `id` (UUID, PK)
- `userId` (FK → User, nullable for guest sessions)

- `role` (from `User.role` or `customer_guest`)
- `channel` (enum: `chat`, `voice`)
- `linkedShipmentId` (FK → Shipment, nullable)
- `openAiSessionId` (string; OpenAI session identifier)
- `startedAt`
- `endedAt` (nullable)
- `status` (enum: `active`, `completed`, `error`)
- `lastAgentName` (e.g., `ShipmentTrackingAgent`)
- `transcript` (JSON or text)
- `outcome` (JSON; e.g., `issueCreated`, `changeRequested`, `escalationTriggered`)

## 9.2 MetricDefinition

- `id` (UUID, PK)
- `key` (unique; e.g., `on_time_delivery_rate`)
- `name`
- `description`
- `aggregationType` (enum: `ratio`, `count`, `avg`)
- `dimension` (enum/string: `global`, `region`, `route`, `driver`)
- `targetValue` (numeric)
- `warningThreshold` (numeric, optional)
- `criticalThreshold` (numeric, optional)
- `ownerRole` (usually `admin`)
- `isVisibleOnDashboard` (bool)
- `createdAt`, `updatedAt`

## 9.3 MetricSnapshot

- `id` (UUID, PK)
- `metricId` (FK → MetricDefinition)
- `value` (numeric)
- `timeRangeStart`
- `timeRangeEnd`
- `computedAt`
- `breakdown` (JSON; e.g., per region)

## 9.4 DashboardConfig

- `id` (UUID, PK)
- `ownerType` (enum: `role`, `user`)
- `ownerRole` (nullable)
- `ownerUserId` (nullable; FK → User)
- `layout` (JSON; widgets + arrangement)
- `createdAt`, `updatedAt`

## 10) WebSocket Events (Minimum Set)

Define and implement events (names are suggestions): - shipment.scan.created - shipment.status.updated - issue.created - issue.updated - escalation.triggered - escalation.advanced - escalation.acknowledged - metrics.snapshot.created (optional)

All events must be: - Authorized (RBAC) - Logged (at least in application logs)

---

## 11) Implementation Plan (Milestones)

### Milestone 1 — Database & Seed Data

- Implement Prisma models + migrations.
- Add seed script:
- 30–50 shipments with scan history
- Routes + route stops
- A few issues + escalation contacts
- Metric definitions + dashboard config defaults

### Milestone 2 — Backend REST + RBAC

- CRUD/queries for shipments, routes, issues.
- Escalation endpoints (trigger, advance, acknowledge).
- Metrics endpoints:
  - GET /metrics/overview
  - GET /metrics/definitions (admin edit)
  - POST /metrics/definitions / PATCH ... (admin only)

### Milestone 3 — WebSockets

- Gateway + event publishing on:
- new scans
- issue updates
- escalation updates
- metric snapshot updates (optional)

### Milestone 4 — OpenAI Agents + Realtime (Chat + Voice)

- Define agents in OpenAI Agents.
- Define tools that map to backend operations.
- Implement backend “AI Orchestrator”:
- Starts/maintains chat sessions
- Starts/maintains Realtime voice sessions
- Logs AgentSession records
- Ensure chat and voice share the same routing + tool calls.

### Milestone 5 — Frontend Dashboards

- Customer dashboard with unified chat + voice widget.
- Admin/Ops dashboard with:

- Shipments, routes, issues, escalations
- Metrics overview + metrics admin editor
- WebSocket subscription handling.

## Milestone 6 — Demo + Documentation

- README with run steps and demo scripts.
  - 3 demo flows: 1) Track shipment → create issue (chat) 2) Request delivery change (voice) 3) SLA risk → escalation → manager ACK (dashboard)
- 

## 12) Deliverables

1. **Backend repository** (NestJS + Prisma + Postgres)
  2. **Frontend repository** (React + TS + Tailwind)
  3. **OpenAI agent configuration** (documented prompts + tools)
  4. **Seed data script** + DB migration files
  5. **README** with setup + demo scripts
- 

## 13) Evaluation Criteria

- Correctness of core flows (tracking, issues, changes, escalations).
- Clean architecture (domain boundaries, services, modules).
- RBAC correctness across REST + WebSockets.
- OpenAI-only compliance (Agents + Realtime for voice; no external AI).
- Professional UX for both customer and admin dashboards.
- Observability: logs + useful error handling.